#### 1 Introduction to OS

OS: intermediary between user and hardware

- time-sharing: illusion of concurrency
- allocate resources, control program
- Kernel; monolithic vs. microkernel has more overhead but smaller
- Hypervisors: T1 (over hardware), T2 (over host OS)

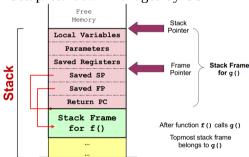
#### 2 Process Abstraction

Process: abstraction of running program

- memory/ hardware/ OS context

# Stack Memory Region

- Stack frame (sf): for function invocation
- stack grows → address decreases
- Setup/ teardown managed by OS



- Frame pointer: fixed location in sf
- Saved registers; GPRs may be temporarily stored in memory first (register spilling)

Caller: Continues execution in caller

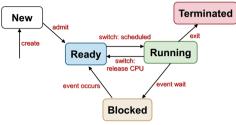
On executing function call: Caller: Pass arguments with registers and/or stack Caller: Save Return PC on stack Transfer control from caller to callee Callee: Save registers used by callee. Save old FP, SF Callee: Allocate space for local variables of callee on stack Callee: Adjust SP to point to new stack top On returning from function call: Callee: Restore saved registers, FP, SP Transfer control from callee to caller using saved PC

### Memory Context

- Heap: dynamically allocated memory (malloc, new)

#### OS Context

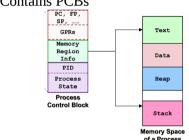
- Process identification (PID) reuse? Limit to maximum no.? Reserved?
- Process state



Hardware Context - GPR, PC, SP, FP

### **Process Table**

- Contains PCBs



- Text: Instructions, Data: Global var
- Pointers are in stack, but memory region pointed to can be in heap

# Interaction with OS

- System Calls: API to OS (involves change to kernel mode, via TRAP). dispatcher finds appropriate system call handler
- Function wrapper (same)/ adapter (modified)
- Exceptions (synchronous, due to program execution) and Interrupts (async. e.g. CTRL-C/ kill)

#### **UNIX Case Study**

- init: root process of process tree
- fork(): creates duplicate process
  - . returns PID of child
  - . can exec() another process
  - . exit(), wait(&status), wait(NULL)
  - . Copy on Write: only duplicate memory when changed
- Zombie processes:
  - . Parent exists before child: init becomes pseudo parent
  - . Child terminates before wait: child becomes zombie (PCB not cleared)

# 3 Process Scheduling

Concurrency/ pseudoparallelism

- Has to be fair, balance of resources
- Scheduler triggered (timer interrupts), decides if context switch is needed
- Cooperative vs. preemptive (fixed time Low priority process can starve (can quota, time slicing)

# **Batch Processing**

- No user interaction, non-preemptive
- Turnaround time: finish arrival time
- Throughput: #tasks/ unit time
- CPU utilisation: %time CPU busy

# First Come First Serve

- FIFO, no starvation
- Convoy effect: long process runs first

# **Shortest Job First**

- Minimises average waiting time
- Starvation possible
- Need to know total CPU time for a task Lottery Scheduling (predicted using exponential average)

# **Shortest Remaining Time**

- New job with shorter remaining time can preempt current running job
- Good service for short jobs, even with late arrivals

#### **Interactive Systems**

- Response time: response request time
- Predictability (variance in response time)
- Preemptive: scheduler runs periodically)
- Interval of timer interrupt (ITI)
- Time quantum: multiple of ITI
- Context switches when time quanta over
- Remaining time quantum: next process can be scheduled

#### **Round Robin**

- FIFO with fixed time slice
- Response time guarantee: (n-1)\*q
- Larger quantum: more CPU utilisation, but longer response time

# **Priority Scheduling**

- Variation: higher priority process can optionally preempt lower priority process
- P1 = highest
- decrease priority after every time quantum)

Priority Inversion: P1 process depends on resource locked by P3, P2 runs instead

# Multi-Level Feedback Oueue (MLFO)

- Scheduling without perfect knowledge, minimises response time for I/O bound processes and turnaround time for CPUbound processes
- Priority reduced if job fully utilises time
  - . Counter abuse using cumulative time . Periodically shift to highest priority

- Can change % of tickets owned
- Responsive: new process has a chance to run

#### **4 Inter-Process Communication**

# **Shared Memory**

- P1 creates M, P2 attached M to its own memspace (synchronisation problems)
- UNIX: shmget, shmat, shmdt, shmctl

# Message Passing

- Msg in kernel memory space
- Direct: explicitly name other party (must know identity)
- Indirect: Mailbox; shared among multiple processes
- Synchronous: blocking

#### **UNIX Pipes**

- Input, output, error
- Circular bounded buffer
- Data accessed in FIFO order
- Producer-Consumer relationship
- dup2 for input/ output redirection

# **UNIX Signals**

- Async (interrupts); kill, stop, continue, errors...

# **5 Threads**

- Lightweight; share code, data and files
- Duplicated registers and stack (SP, FP points to different location in same main stack)
- More resource sharing, responsive, scalable, less overhead, less protection
- Can execute different threads in parallel
- fork(): usually only one thread. Exit()? Exec()? Which thread handles signal?
- User thread: library, kernel unaware
   One thread blocks → whole process blocks
- Kernel thread: can have thread-level scheduling

- POSIX: pthread (-lpthread)
  - . Shared memory space
  - . pthread\_join to synchronise

# **6 Synchronisation**

### **Race Condition**

- Non-atomic instructions, outcome depends on order of execution

### Critical Section (CS)

- Mutual exclusion, progress, bounded wait, independence *Problems* 

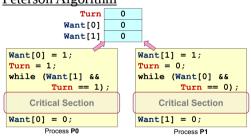
- Deadlock: all blocked
- Live Lock: no progress due to deadlock avoidance mechanism
- Starvation: some processes blocked forever

#### TestAndSet

- Turn:

- Atomic, assembly-level
- Load content at memory location into register, set memlocation = 1
- Busy waiting (processes still scheduled)
- Bounded wait: depends on scheduler

# Peterson Algorithm



- Want; ensures independence, P1 does not have to wait for P0 to enter CS first
  - . Ensures mutual exclusion: only holds one value at each time

- . Prevents deadlock: both "want" = 1
- Busy waiting
- Only synchronises 2 processes

# **Semaphore**

- Atomic wait, signal
- Wait: s <= 0: block; s--;
- Signal: s++; wakes up one sleeping process (if any)
- Invariant: S<sub>current</sub> = S<sub>initial</sub> + #signal(s) -#wait(s), wait() must be completed
- Binary semaphore: mutex
- #in CS = #wait(s) #signal(s) <= 1
- No deadlock ( $S_{\text{current}}$  and  $N_{\text{CS}}$  cannot both be 0)
  - . Unless two semaphores blocked
- No starvation, assuming fair scheduling

Producer-Consumer (e.g pipe)

```
while (TRUE) {
      Produce Item;
                                       wait( notEmpty );
      wait( notFull );
                                       wait( mutex ):
     wait( mutex );
buffer[in] = item;
                                       item = buffer[out];
                                       out = (out+1) % K;
      in = (in+1) % K;
      count++;
                                       signal( mutex );
      signal ( mutex );
                                       signal ( notFull );
      signal ( notEmpty )
                                       Consume Item:
             Producer Process
                                              Consumer Process
```

# **Dining Philosophers**

- deadlock: all have left only
- livelock: all unable to pick right
- States: Think, Hungry, Eat

#### Solution 1: Tanenbaum

```
void takeChpStcks( i )
{
    wait( mutex );
    state[i] = HUNGRY;
    safeToEat( i );
    signal( mutex );
    wait( s[i] );
}
```

```
void putChpStcks( i )
{
    wait( mutex );
    state[i] = THINKING;
    safeToEat( LEFT );
    safeToEat( RIGHT );
    signal( mutex );
}
```

- signals to left and right that chopsticks are ready

# Reader-Writer (e.g. files)

```
while (TRUE)
while (TRUE) {
                                   wait( mutex );
   wait( roomEmpty );
                                   nReader++;
                                   if (nReader == 1)
                                       wait( roomEmpty );
    Modifies data
                                   signal( mutex );
   signal( roomEmpty );
                Writer Process
                                   wait( mutex )
                                   nReader--;
                                  if (nReader == 0)
                                       signal( roomEmpty );
 Initial Values:
                                   signal ( mutex );
 o roomEmpty = S(1)
 mutex = S(1)
                                               Reader Process
 n nReader = 0
```

- mutex used to protect CS of nReader
- writer can be starved

# Solution 2:

```
void philosopher( int i ) {
    while (TRUE) {
        Think( );
        wait( seats );
        wait( chpStk[LEFT] );
        wait( chpStk[RIGHT] );
        Eat( );
        signal( chpStk[LEFT] );
        signal( chpStk[RIGHT] );
        signal( seats );
    }
}
```

- Limited Eater: n -1 people for n seats